

DATA SCIENCE CAPSTONE PROJECT

MACHINE LEARNING GPU CLUSTERING

Implementation

Setting up Kubernetes

STEP 1

Disabling the firewall

```
ufw disable
```

STEP 2

Disabling Swap

```
swapoff -a  
sed -i '/swap/d' /etc/fstab
```

STEP 3

Update sysctl settings for Kubernetes networking

```
cat >>/etc/sysctl.d/kubernetes.conf<<EOF  
net.bridge.bridge-nf-call-ip6tables = 1  
net.bridge.bridge-nf-call-iptables = 1  
EOF  
sysctl --system
```

STEP 4

Install docker engine

```
apt install -y apt-transport-https ca-certificates curl gnupg-agent software-properties-common  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | apt-key add -  
add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -  
cs) stable"  
apt update
```

```
apt install -y docker-ce=5:19.03.10~3-0~ubuntu-focal containerd.io
```

STEP 5

Kubernetes Setup

Add Apt repository

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -  
echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" > /etc/apt/sources.list.d/kubernetes.list
```

STEP 6

Install Kubernetes components

```
apt update  
apt install -y kubeadm=1.18.5-00 kubelet=1.18.5-00 kubectl=1.18.5-00
```

ONLY ON K-MASTER

Initialize Kubernetes Cluster

Update the below command with the ip address of kmaster

```
kubeadm init --apiserver-advertise-address=172.16.16.100 --pod-network-cidr=192.168.0.0/16 --  
ignore-preflight-errors=all
```

Deploy Calico network

```
kubectl --kubeconfig=/etc/kubernetes/admin.conf create -f  
https://docs.projectcalico.org/v3.14/manifests/calico.yaml
```

Cluster join command

```
kubeadm token create --print-join-command
```

To be able to run kubectl commands as non-root user

If you want to be able to run kubectl commands as non-root user, then as a non-root user perform these

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

On Kworker

Join the cluster

Use the output from **kubeadm token create** command in previous step from the master server and run here.

Use kubectl get nodes command to check if the nodes are ready.

Installing helm

Run the helm installer.

```
curl https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3 | bash
```

Verify helm.

```
helm list
```

Setting up NFS Subdir External Provisioner

Install NFS Kernel Server in Ubuntu

```
sudo apt update
```

```
sudo apt install nfs-kernel-server
```

Create NFS Export Directory

```
sudo mkdir -p /srv/nfs/kubedata
```

```
sudo chown -R nobody:nogroup /srv/nfs/kubedata
```

Edit the etc exports file

```
sudo vi /etc/exports
```

add the following line

```
/srv/nfs/kubedata *(rw,sync,no_subtree_check,no_root_squash,no_all_squash,insecure)
```

Save the file

Enable and start the nfs server

```
sudo systemctl enable --now nfs-server
```

```
sudo exportfs -rav
```

Run `sudo showmount -e localhost` it tells you what directories are being exported through the nfs

To verify login to your Kubernetes nodes check if the nodes can mount nfs volumes

```
apt install -y nfs-common
```

```
mount -t nfs ip_address:/srv/nfs/kubedata
```

check the mount and unmount it

```
umount /mnt
```

Check if the system already have nfs server if not configure nfs server

Helm install NFS subdir external provisioner

First add helm repo

```
helm repo add nfs-subdir-external-provisioner https://kubernetes-sigs.github.io/nfs-subdir-external-provisioner/
```

Install NFS subdir external provisioner

```
helm install nfs-subdir-external-provisioner nfs-subdir-external-provisioner/nfs-subdir-external-provisioner \ --set nfs.server=x.x.x.x \ --set nfs.path=/exported/path
```

Run `kubectl get pods` to check nfs pod is running

Set up Metallb

We have to enable strict ARP mode

```
kubectl edit configmap -n kube-system kube-proxy
```

set `strictARP` as true

```
apiVersion: kubeproxy.config.k8s.io/v1alpha1
```

```
kind: KubeProxyConfiguration
```

```
mode: "ipvs"
```

```
ipvs:
```

```
  strictARP: true
```

Run the following commands

```
kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.9.6/manifests/namespace.yaml
```

```
kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.9.6/manifests/metallb.yaml
```

Create metal_config.yaml using the following file

```
vi metal_config.yaml
```

Replace PUBLIC_IP_ADDRESS with the actual public ip address of your instance

```
apiVersion: v1
```

```
kind: ConfigMap
```

```
metadata:
```

```
  namespace: metallb-system
```

```
  name: config
```

```
data:
```

```
  config: |
```

```
    address-pools:
```

```
    - name: default
```

```
      protocol: layer2
```

```
      addresses:
```

```
      - PUBLIC_IP_ADDRESS-PUBLIC_IP_ADDRESS
```

Apply custom metal load balancer configuration — metal_config.yaml

```
kubectl apply -f metal_config.yaml
```

Check the metal load balancer status

```
kubectl --namespace=metallb-system get pods
```

Setting up Jupyterhub

Prepare configuration file

1. Generate a random hex string representing 32 bytes to use as a security token. Run this command in a terminal and copy the output:

```
openssl rand -hex 32
```

2. Create and start editing a file called config.yaml. In the code snippet below we start the widely available nano editor

```
nano config.yaml
```

3. Write the following into the config.yaml file but instead of writing <RANDOM-HEX> paste the generated hex string you copied in step 1.

proxy:

```
secretToken: "<RANDOM-HEX>"
```

hub:

db:

```
type: sqlite-memory
```

singleuser:

storage:

dynamic:

```
storageClass: nfs-client
```

1. Save the config.yaml file. In the nano editor this is done by pressing **CTRL+X** or **CMD+X** followed by a confirmation to save the changes.

Install JupyterHub

1. Make Helm aware of the JupyterHub Helm chart repository so you can install the JupyterHub chart from it without having to use a long URL name.

```
helm repo add jupyterhub https://jupyterhub.github.io/helm-chart/helm repo update
```

2. Now install the chart configured by your config.yaml by running this command from the directory that contains your config.yaml:

```
RELEASE=jhub
```

```
NAMESPACE=jhub
```

```
helm upgrade --cleanup-on-fail \ --install $RELEASE jupyterhub/jupyterhub \ --namespace $NAMESPACE \ --create-namespace \ --version=0.10.6 \ --values config.yaml
```

3. While Step 2 is running, you can see the pods being created by entering in a different terminal:

```
kubectl get pod --namespace jhub
```

4. To remain sane we recommend that you enable autocompletion for kubectl and set a default value for the --namespace flag:

```
kubectl config set-context $(kubectl config current-context) --namespace ${NAMESPACE:-jhub}
```

5. Find the IP we can use to access the JupyterHub. Run the following command until the EXTERNAL-IP of the proxy-public service is available like in the example output.

```
Kubectl get service --namespace jhub
```

Your jupyterhub setup is ready

References

<https://github.com/justmeandopensource/kubernetes/blob/master/docs/install-cluster-ubuntu-20.md>

<https://github.com/kubernetes-sigs/nfs-subdir-external-provisioner>

<https://zero-to-jupyterhub.readthedocs.io/en/stable/jupyterhub/installation.html>

<https://georgepaw.medium.com/jupyterhub-with-kubernetes-on-single-bare-metal-instance-tutorial-67cbd5ec0b00>

<https://metallb.universe.tf/installation/>

<https://dev.to/upindersujlana/upgrade-kubernetes-cluster-to-1-19-4-using-kubeadm-3ien>

<https://www.tecmint.com/install-nfs-server-on-ubuntu/>

<https://www.youtube.com/watch?v=DF3v2P8ENeg&t=188s>